



# **MPI**

***MESSAGE PASSING INTERFACE***

# NON-BLOCKING FUNCTIONS

- Previous functions are blocking: the following instruction is not executed until they have finished.
- Non-blocking functions: function returns, whether data transfer is finished or not.
  - Requires function to query the status of the data transfer
  - Message buffers are needed (length of message is limited)
  - Overlapping of communication and computation is possible (reduction of execution time)
- Examples:
  - MPI\_Isend
  - MPI\_Irecv
  - MPI\_Iprobe
  - etc..

# NON-BLOCKING FUNCTIONS

- `MPI_Isend`: sends a message but returns before copying into the buffer (buf cannot be modified before data is received).
  - `int MPI_Isend(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_COMM comm, MPI_Request *request)`
    - `buf`: buffer with data to be sent
    - `count`: size of data to be sent
    - `datatype`: type of data to be sent
    - `dest`: destiny process id (where data are sent to)
    - `tag`: message id
    - `comm`: communicator
    - `request`: to identify the operation in progress.

# NON-BLOCKING FUNCTIONS

- `MPI_Irecv`: receives a message but returns before copying into the buffer. Two possibilities: waiting till it ends (`MPI_Wait`) or checking the communication's status (`MPI_Test`).

- `int MPI_Irecv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_COMM comm, MPI_Request *request)`

- `buf`: buffer where received data will be stored
- `count`: size of data to be received
- `datatype`: type of data to be received
- `source`: source process id (where data come from)
- `tag`: message id
- `comm`: communicator
- `request`: to identify the operation in progress.

# NON-BLOCKING FUNCTIONS

- `MPI_Test`: checks if the non-blocking operation has finalized.
  - `int MPI_Test(MPI_Request *request, int *flag, MPI_Status *status)`
    - `request`: to identify the operation we are waiting for.
    - `status`: gathers information about the non-blocking operation that has finalized (pointer)
    - If `flag=true` then the operation has finalized; `request` is freed, `status` is initialized

# NON-BLOCKING FUNCTIONS

- `MPI_Wait`: blocks till the non-blocking operation has finished.
  - `int MPI_Wait(MPI_Request *request, MPI_Status *status)`
    - `request`: to identify the operation we are waiting for.
    - `status`: gathers information about the non-blocking operation that has finalized (pointer)

# NON-BLOCKING FUNCTIONS

- `MPI_Iprobe`: checks if there are messages to receive.
  - `int MPI_IProbe(int source, int tag, MPI_Comm comm, int *flag, MPI_Status *status)`
    - `source`: source process id (where we check if it has sent data)
    - `tag`: message id (the one that we check if it has been received)
    - `comm`: communicator
    - `status`: gathers information about the message to be received (pointer)
    - If `flag=true` the message fits (`source, tag, comm`). Status gives more info.
    - `MPI_ANY_SOURCE` and `MPI_ANY_TAG` allow us to wait for messages from any source and/or with any tag.
    - It does not block if there are no messages. If there are messages, they are received with `MPI_Recv`.

# NON-BLOCKING FUNCTIONS

- `MPI_Probe`: checks for specific messages.
  - `int MPI_Probe(int source, int tag, MPI_Comm comm, MPI_Status *status)`
    - `source`: source process id (where we check if it has sent data)
    - `tag`: message id (the one that we check if it has been received)
    - `comm`: communicator
    - `status`: gathers information about the message to be received (pointer)
    - `MPI_ANY_SOURCE` and `MPI_ANY_TAG` allow us to wait for messages from any source and/or with any tag.
    - It blocks: it finishes when there are messages that fit the given arguments. If there are messages, they are received with `MPI_Recv`.



# EXAMPLE

```
...
int datasize, *buf, source;
...
//Comprobamos si hay algún mensaje pendiente
MPI_Probe(MPI_ANY_SOURCE, 0, comm, &status);
//Obtenemos el tamaño de los datos a recibir y
//reservamos memoria
MPI_Get_count(status, MPI_INT, &datasize);
buf=malloc(datasize*sizeof(int));
//Obtenemos el proceso fuente del mensaje
source=status.MPI_SOURCE;
//Recibimos el mensaje
MPI_Recv(buf,datasize,MPI_INT,source,0,comm,&status);
...
```

# NON-BLOCKING FUNCTIONS

- Waits for any specified send or receive to complete:
  - `int MPI_Waitany (int count, MPI_Request *array_of_requests, int *index, MPI_Status *status);`
    - `count`: list length
    - `array_of_requests`: array of requests
    - `index`: index of handle for operation that completed
    - `status`: status object
  - It is used to wait for the completion of one out of several requests.

# NON-BLOCKING FUNCTIONS

- Waits for all given communications to complete:
  - `int MPI_Waitall (int count,  
MPI_Request *array_of_requests,  
MPI_Status *array_of_statuses);`
    - count: list length
    - array\_of\_requests: array of requests
    - array\_of\_statuses: array of status objects
  - Blocks until all communication operations associated with active handles in the list complete, and returns the status of all these operations.
    - This includes the case where no handle in the list is active

# NON-BLOCKING FUNCTIONS

- Waits for some given communications to complete:
  - `int MPI_Waitsome (int incount, MPI_Request *array_of_requests, int *outcount, int *array_of_indices, MPI_Status *array_of_statuses);`
    - `incount`: length of `array_of_requests`
    - `array_of_requests`: array of requests
    - `outcount`: number of completed requestes
    - `array_of_indices`: array of indices of operations to be completed
    - `array_of_statuses`: array of status objects for operations that completed
  - Waits until at least one of the operations associated with active handles in the list have completed.

# NON-BLOCKING FUNCTIONS

- Tests for completion of any one previously initiated communication in a list:
  - `int MPI_Testany (int count, MPI_Request *array_of_requests, int *index, int *flag, MPI_Status *status);`
    - `count`: list length
    - `array_of_requests`: array of requests
    - `index`: index of operation that completed or `MPI_UNDEFINED` if none completed
    - `flag`: true if one of the operations is completed
    - `status`: status object
  - It tests for completion of either one or none of the operations associated with active handles.

# NON-BLOCKING FUNCTIONS

– Tests for the completion of all previously initiated communications in a list:

- `int MPI_Testall (int count, MPI_Request *array_of_requests, int *flag, MPI_Status *array_of_statuses);`

- `count`: list length

- `array_of_requests`: array of requests

- `flag`: true if previously initiated communications are complete

- `array_of_statuses`: array of status objects

- Returns `flag=true` if all communications associated with active handles in the array have complete (this includes the case where no handle in the list is active)

# NON-BLOCKING FUNCTIONS

– Tests for completion of one or more previously initiated communications in a list:

- `int MPI_Testsome (int incount, MPI_Request *array_of_requests, int *outcount, int *array_of_indices, MPI_Status *array_of_statuses);`
  - `incount`: length of `array_of_requests`
  - `array_of_requests`: array of requests
  - `outcount`: number of completed requestes
  - `array_of_indices`: array of indices of operations that completed
  - `array_of_statuses`: array of status objects for operations that completed
- Behaves like `MPI_Waitsome` except that it returns immediately

# AVOIDING INTERLOCKS

- An interlock is produced:
  - When one or more processes achieves a blocking receiving routine but the message never comes. Process waits indefinitely and there is no error message.
  - E.g. 2 processes interchange messages and we do not program it well.



# EXAMPLE: IT ALWAYS WORKS

...

```
if (myid == 0) {  
    MPI_Send(&a, 1, MPI_FLOAT, 1, tag, MPI_COMM_WORLD);  
    MPI_Recv(&b, 1, MPI_FLOAT, 1, tag, MPI_COMM_WORLD, &status);  
} elseif (myid == 1) {  
    MPI_Recv(&a, 1, MPI_FLOAT, 0, tag, MPI_COMM_WORLD, &status);  
    MPI_Send (&b, 1, MPI_FLOAT, 0, tag, MPI_COMM_WORLD);  
}
```

...

# EXAMPLE: IT NEVER WORKS

...

```
if (myid == 0) {  
    MPI_Recv (&b, 1, MPI_FLOAT, 1, tag, MPI_COMM_WORLD, &status)  
    MPI_Send(&a, 1, MPI_FLOAT, 1, tag, MPI_COMM_WORLD);  
} elseif (myid == 1) {  
    MPI_Recv(&a, 1, MPI_FLOAT, 0, tag, MPI_COMM_WORLD, &status);  
    MPI_Send(&b, 1, MPI_FLOAT, 0, tag, MPI_COMM_WORLD);  
}
```

...

# EXAMPLE: IT MAY WORK

...

```
if (myid == 0) {  
    MPI_Send(&a, 1, MPI_FLOAT, 1, tag, MPI_COMM_WORLD);  
    MPI_Recv(&b, 1, MPI_FLOAT, 1, tag, MPI_COMM_WORLD, &status)  
} elseif (myid == 1) {  
    MPI_Send(&b, 1, MPI_FLOAT, 0, tag, MPI_COMM_WORLD);  
    MPI_Recv(&a, 1, MPI_FLOAT, 0, tag, MPI_COMM_WORLD, &status);  
}
```

...



Avoid it!

# EXAMPLE: IT MAY WORK

...

```
if (myid == 0) {  
    MPI_Send(&a, 1, MPI_FLOAT, 1, tag, MPI_COMM_WORLD);  
    MPI_Recv(&b, 1, MPI_FLOAT, 1, tag, MPI_COMM_WORLD, &status)  
} elseif (myid == 1) {  
    MPI_Send(&b, 1, MPI_FLOAT, 0, tag, MPI_COMM_WORLD);  
    MPI_Recv(&a, 1, MPI_FLOAT, 0, tag, MPI_COMM_WORLD, &status);  
}
```

...

Recv() ALWAYS blocks until the corresponding Send() is carried out

Depending on the platform, Send() MAY block until the corresponding Recv() is carried out

# AVOIDING INTERLOCKS

– MPI\_Sendrecv: sending and receiving data at the same time.

- `int MPI_Sendrecv(void *sendbuf, int sendcount, MPI_Datatype sendtype, int dest, int sendtag, void *recvbuf, int recvcount, MPI_Datatype recvtype, int source, int recvtag, MPI_Comm comm, MPI_Status *status)`

`sendbuf`: **buffer** of data to be sent

`sendcount`: size of data to be sent // `sendtype`: type of data to be sent

`dest`: destiny process id // `sendtag`: tag of the sending data

`recvbuf`: **buffer** of data to be received

`recvcount`: size of data to be received // `recvtype`: type of data to be received

`source`: source process id // `recvtag`: tag of the receiving data

`comm`: communicator

`status`: gathers information about the receiving operation (pointer)

# EXAMPLE

...

```
tag1=1;
```

```
tag2=2;
```

```
if (myid == 0) {
```

```
    MPI_Sendrecv(&a,1, MPI_FLOAT,1,tag1, &b,1,  
    MPI_FLOAT,1,tag2, MPI_COMM_WORLD,&status);
```

```
elseif (myid == 1) {
```

```
    MPI_Sendrecv(&b,1, MPI_FLOAT,0,tag2, &a,1,  
    MPI_FLOAT,0,tag1, MPI_COMM_WORLD,&status);
```

```
}
```

...



Recomended!

# AVOIDING INTERLOCKS

– `MPI_Sendrecv_replace`: performs send and receive in one single function call and operates only one single buffer.

- `int MPI_Sendrecv_replace(void *buf, int count, MPI_Datatype type, int dest, int sendtag, int source, int recvtag, MPI_Comm comm, MPI_Status *status)`

`buf`: **buffer** of data to be sent/to store the received data

`count`: **size** of message (in elements)

`type`: **type** of data

`dest`: **destiny** process id    // `sendtag`: **tag** of the sending data

`source`: **source** process id // `recvtag`: **tag** of the receiving data

`comm`: **communicator**

`status`: **gathers** information about the receiving operation (pointer)

# EXERCISE

Write a MPI program to communicate  $N$  processes in a ring way.  $i$  Process must tell  $i+1$  process its machine name.

- Implementation 1: use `send()` and `recv()`.
- Implementation 2: use `Isend()` and `Irecv()`.
- Implementation 3: use `Sendrecv()`.
- In implementations 1 and 2, `Send()/Isend()` and `Recv()/Irecv()` can be executed in any order?